Jordan Taylor Chen

jordan.chen@example.com | +1-555-555-5555 | linkedin.com/in/jordanchen | github.com/jordanchen

PROFESSIONAL SUMMARY

Senior AI Engineer with 6+ years of experience building cloud-native platforms and production LLM systems. Currently leading development of an enterprise retrieval and evaluation platform at Acme Corp, orchestrating RAG pipelines across multiple model providers. Proven record of shipping applied AI from 0→1 and scaling it to thousands of users without drifting on quality or cost.

WORK EXPERIENCE

Acme Corp **Remote**

Senior AI Engineer, Applied AI Platform *04/2024 – present*

- Architected and deployed a production RAG platform integrating a primary LLM provider, a vector database, and a structured-output grounding layer within a cloud-native architecture, serving 500+ enterprise workspaces since public launch.
- Led the evaluation pipeline design using golden-set scoring plus online sampling, cutting regression-detection time 30% versus the prior manual spot-check process.
- Produced 25+ internal architecture guides and reference implementations that reduced platform onboarding support tickets by 22%.
- Led cross-functional collaboration across product and engineering to translate production feedback into platform improvements, lifting answer-quality scores 12%.

Beacon Systems **Remote**

Platform Engineer, Partnerships *01/2023 – 03/2024*

- Built partner integrations using Ruby on Rails, React, and TypeScript, automating service workflows across the platform.
- Developed integration frameworks that increased partner engagement by 45% and reduced integration time by 2 weeks.
- Led coordination between Product and Engineering to convert platform feedback into API improvements, contributing to a 20% increase in overall platform adoption.

Northwind Consulting **San Francisco, CA**

Software Engineer *07/2021 – 12/2022*

- Modernized enterprise applications by migrating legacy systems to Rails-based infrastructure, decreasing delivery time by 15%.
- Facilitated architecture workshops with engineering teams and business stakeholders to translate complex workflows into scalable systems.
- Authored technical documentation that reduced integration time by 30% and improved onboarding consistency.

TECHNICAL PROJECTS

Agent Playground: Open-Source LLM Evaluation Harness

github.com/jordanchen/agent-playground

- Built open-source eval harness for multi-turn LLM agents with golden-set scoring, online sampling, and regression detection across model upgrades.
- Shipped reference implementations for five common agent patterns, adopted by 200+ developers since release.

Chatgrounding: RAG Reference Implementation

github.com/jordanchen/chatgrounding

- Built full-stack RAG reference on a public LLM API with document grounding and multi-turn memory, demonstrating applied RAG without heavy frameworks.
- Published as an open-source starter for businesses deploying embedded AI chat: marketing site, chat widget, and retrieval layer.

VOLUNTEER COMMUNITY INVOLVEMENT

AI Builders Meetup: Organizer *2022 – present*

- Organized monthly technical meetups on LLM systems engineering, growing regional attendance to 300+ recurring members.
- Hosted educational sessions on applied AI integration patterns, lifting session engagement 90% year over year.

SKILLS

- **AI Systems:** OpenAI / Anthropic / Cohere APIs, LLM integration, agent orchestration, RAG, multi-model evaluation, context engineering
- **Data + retrieval:** vector databases, embeddings, hybrid search, document chunking strategies

- **API design:** REST + streaming APIs, documentation, integration architecture
- **DevOps:** Docker, CI/CD, GitHub Actions, observability
- **Languages:** Python, TypeScript, Ruby, JavaScript
- **Frameworks:** Ruby on Rails, React, Node.js

EDUCATION

- **State University:** B.S., Computer Science
- **Online AI Program:** Certificate, Applied LLM Systems